

JASON BAK

jasonbak.co | jason.j.bak@gmail.com | Seattle, WA

EXPERIENCE

HubSpot | Senior Software Engineer
Infra Costs

2020 - 2025
Seattle, WA

*Worked within Platform Infrastructure, building from scratch AWS cost optimization and efficiency systems. Achieved significant cost savings at a company with **\$210M** in annual cloud spend.*

- Identified waste patterns, advocated for prioritization, and implemented an automated resource optimization system across high-cost resource types, **eliminating \$2M in accrued waste** with **zero operational disruptions** spanning all organizations.
- Built a multi-layered automated enterprise cost allocation tagging system (event-driven workers, scheduled remediation jobs, creation-time injection), achieving **\$208M annual cost coverage**. This is the foundation of our cost allocation datasets, alerting infrastructure, and optimization tools.
- Architected configurable cost anomaly alerting system with custom grouping dimensions, multiple threshold types (percentage, scalar, statistical), alert muting, and team-based routing, **preventing \$1M in runaway costs** and **mentoring teammates** on implementation.
- Built a centralized AWS resource metadata service, **tracking 2M resource states**, enabling historical state queries for cost data reprocessing and simplifying data retrieval across internal sources and AWS.
- **Improved annual cost attribution accuracy by \$56M** through a Spark ETL data pipeline that processed AWS billing records. Changes include custom CPU/GPU allocation for shared Kubernetes workloads, proportional EC2 storage cost calculations, and EBS event processors.
- Collaborated directly with infrastructure and product teams to resolve their cost spend issues (waste, inefficiency, runaway costs, "unexplained" costs), **reducing costs by \$2M**.

Intuit | Software Engineer Intern
Data Platforms

2019
Mountain View, CA

- Automated EMR cluster rotation workflow for an internal data infrastructure console, **reducing an 11-minute manual process to one click**, through orchestrated termination, polling, and recreation.

Amazon | Software Development Engineer Intern
AI Algorithms

2018
Seattle, WA

- Enhanced anomaly detection by capturing temporal context in time-series data for the Random Cut Forest C++ library. Included Python and Java bindings for cross-language interoperability.

EDUCATION

Carnegie Mellon University | B.S. in Statistics and Machine Learning
Minor in Computer Science

2019
Pittsburgh, PA

Teaching Assistant: Intro to Machine Learning, Data Structures (2x)

Research Assistant: Institute for Software Research (Machine Learning)

SKILLS

Languages: Java, Python, C++, C, R, JavaScript

Technologies: Spark, MySQL, Kafka, Elasticsearch, AWS (SQS, EventBridge, Athena, S3, DynamoDB), Maven, Grafana, Prometheus, Flask, Snowflake, Git